

Name: Benoit Ortalo-Magne, Leon Medvinsky, Steven Lee
NetId: beo2, leonkm2, solee2
Team name: ChampaignWithoutTheCham
School: Illinois

Milestone 2

```
* Running /bin/bash -c "./m2 10000"
Test batch size: 10000
Loading fashion-mnist data...Done
Loading model...Done
Conv-CPU==
Op Time: 89215.8 ms
Conv-CPU==
Op Time: 263190 ms
Test Accuracy: 0.8714
```

Test Accuracy: .8714
Op Times: 83972.5 ms and 243484 ms
Total Time: 327456.5 ms

Milestone 3

Rai running your GPU implementation of convolution:

```
* Running bin/bash -c "./m3 10000"
Test batch size: 10000
Loading fashion-mnist data...Done
Loading model...Done
Conv-GPU==
Conv-GPU==

Test Accuracy: 0.8714
-----
-              TIMINGS
-----
Layer 1 GPUPTime: 38.464247 ms
Layer 1 OpTime: 38.473143 ms
Layer 1 LayerTime: 631.439901 ms
Layer 2 GPUPTime: 151.761218 ms
Layer 2 OpTime: 151.77021 ms
Layer 2 LayerTime: 592.337754 ms
```

Nsys profiling GPU execution:

```

Generating CUDA API Statistics...
CUDA API Statistics (nanoseconds)

Time(%)    Total Time    Calls    Average    Minimum    Maximum    Name
-----
85.9      1118091706      6      186348617.7    85329    535389563    cudaMemcpy
13.7      178500915      6      29750152.5    64024    176991203    cudaMalloc
0.3       4431407      6      738567.8    58046    2849717    cudaFree
0.0       262923      2      131461.5    25790    237133    cudaLaunchKernel

Generating CUDA Kernel Statistics...

Generating CUDA Memory Operation Statistics...
CUDA Kernel Statistics (nanoseconds)

Time(%)    Total Time    Instances    Average    Minimum    Maximum    Name
-----
100.0     138749857      2     69374928.5    16497215    122252642    conv_forward_kernel

CUDA Memory Operation Statistics (nanoseconds)

Time(%)    Total Time    Operations    Average    Minimum    Maximum    Name
-----
92.6     895095099      2     447547549.5    376998837    518096262    [CUDA memcpy DtoH]
7.4      71795884      4     17948971.0    1248    38466642    [CUDA memcpy HtoD]

CUDA Memory Operation Statistics (KiB)

Total    Operations    Average    Minimum    Maximum    Name
-----
1722500.0      2     861250.0    722500.000    1000000.0    [CUDA memcpy DtoH]
538919.0      4     134729.0      0.766    288906.0    [CUDA memcpy HtoD]

Generating Operating System Runtime API Statistics...
Operating System Runtime API Statistics (nanoseconds)

Time(%)    Total Time    Calls    Average    Minimum    Maximum    Name
-----
33.3     94840706872     963     98484638.5    51075    100207899    sem_timedwait
33.3     94788430779     961     98635203.7    57366    100773798    poll
22.1     62788851699      2    31394425849.5    22305165637    40483686062    pthread_cond_wait
11.2     32009136178     64     500142752.8    500084730    500210532    pthread_cond_timedwait
0.0     101001684     764     132201.2    1063    15635483    ioctl
0.0     19053488     9072     2100.3    1220    18841    read
0.0     2712393      97     27962.8    1043    970934    mmap
0.0     627899      97     6473.2    1716    21485    open64
0.0     625833      2     312916.5    40823    585010    pthread_mutex_lock
0.0     220436      5     44087.2    31233    53756    pthread_create

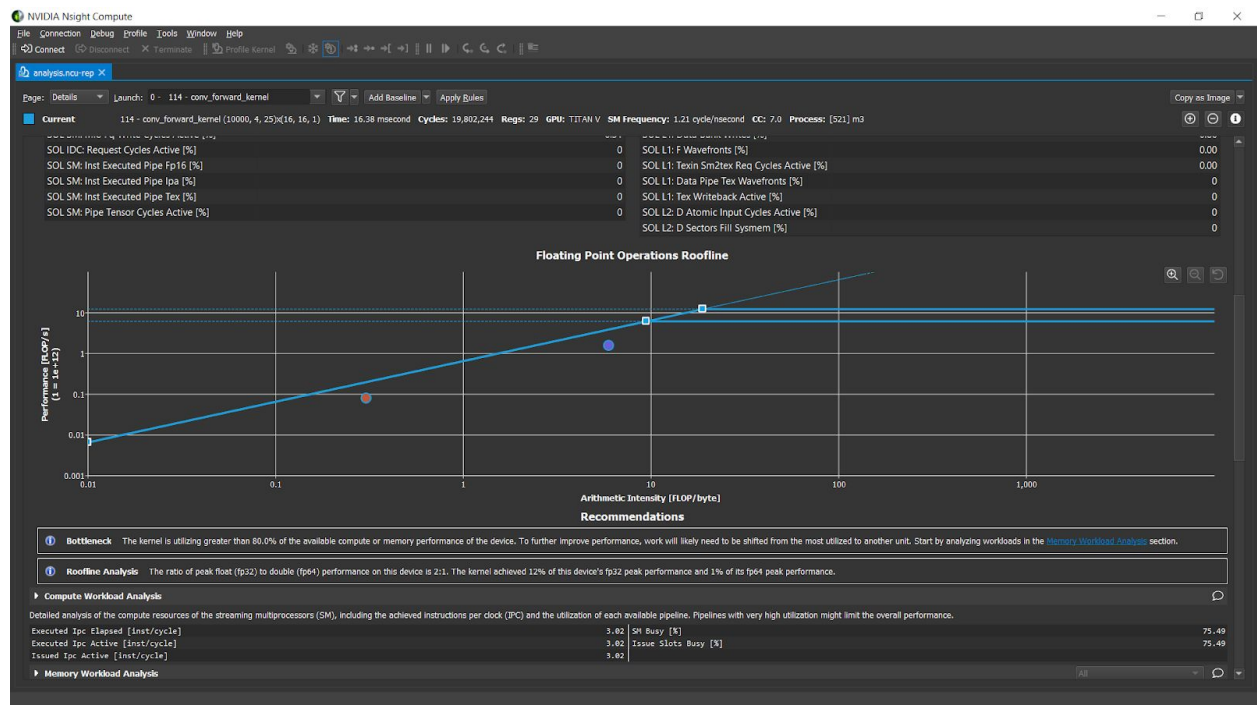
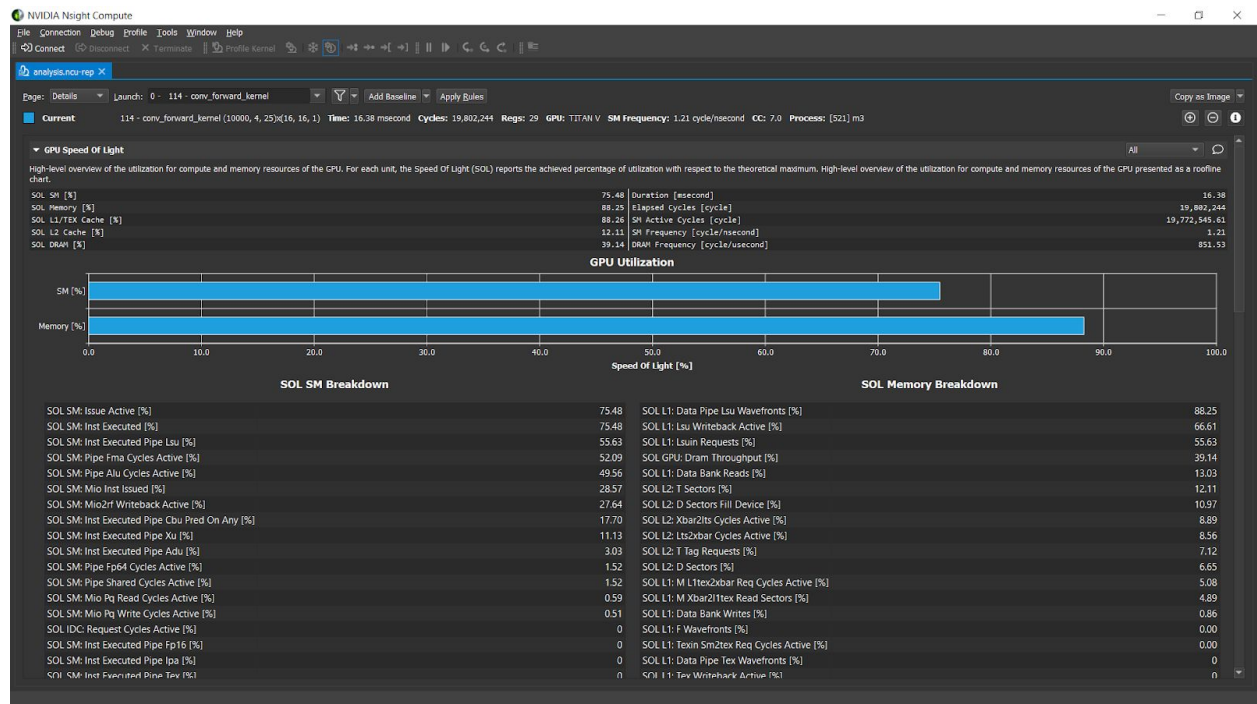
```

Kernels that consume more than 90% of program time: conv_forward_kernel

CUDA API calls that consume more than 90% of program time: cudaMemcpy, cudaMalloc, cudaLaunchKernel, cudaFree

Difference between kernels and API calls: API calls are used to set up memory and variables for kernels to run on.

Screenshot of GPU SOL utilization:



Milestone 4

1. Shared memory

We noticed that the baseline kernel has high memory bandwidth utilization compared to SM utilization, suggesting the kernel was memory-bound. Since convolution (especially batched in 3d) involves a lot of memory reuse between threads, using shared memory to reduce the number of global memory reads was an obvious way to improve the kernel's memory bandwidth requirements.

Effect

Overall timing information is as follows:

Layer 1 GPUTime: 40.26357 ms

Layer 1 OpTime: 40.302002 ms

Layer 1 LayerTime: 638.512441 ms

Layer 2 GPUTime: 226.928337 ms

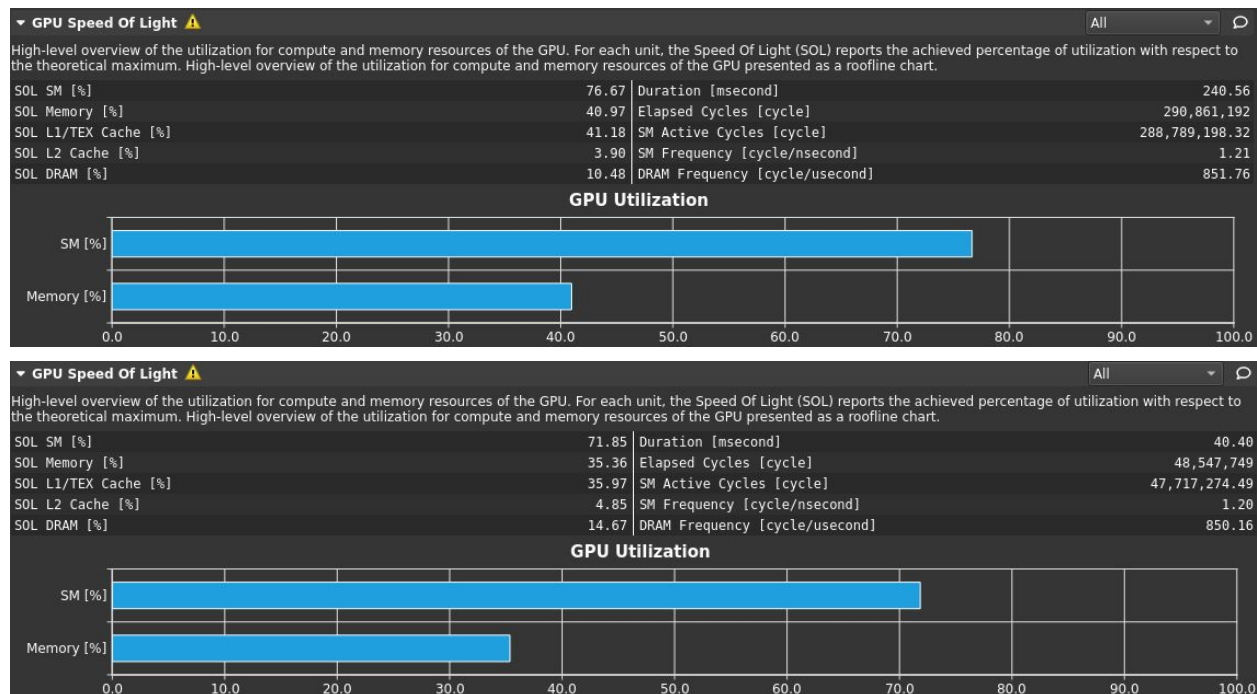
Layer 2 OpTime: 226.961872 ms

Layer 2 LayerTime: 681.999028 ms

Compared to the baseline, performance is considerably worse. Nsys does not shed much light on this, mainly showing the increased execution time.

Time(%)	Total Time	Calls	Average	Minimum	Maximum	Name
81.4	1317579041	8	164697380.1	15431	628180643	cudaMemcpy
17.4	282160770	8	35270096.2	74170	272446462	cudaMalloc
1.0	15821774	6	2636962.3	18114	15696648	cudaLaunchKernel
0.2	2767247	8	345905.9	72194	975341	cudaFree
0.0	22261	4	5565.2	2393	12987	cudaDeviceSynchronize
Generating CUDA Kernel Statistics...						
Generating CUDA Memory Operation Statistics...						
CUDA Kernel Statistics (nanoseconds)						
Time(%)	Total Time	Instances	Average	Minimum	Maximum	Name
100.0	277074773	2	138537386.5	40595080	236479693	conv_forward_kernel
0.0	2688	2	1344.0	1312	1376	do_not_remove_this_kernel
0.0	2592	2	1296.0	1248	1344	prefn_marker_kernel
CUDA Memory Operation Statistics (nanoseconds)						
Time(%)	Total Time	Operations	Average	Minimum	Maximum	Name
92.8	960835387	2	480417693.5	390890544	569944843	[CUDA memcpy DtoH]
7.2	74126576	6	12354429.3	1216	38931280	[CUDA memcpy HtoD]

The output of Nsight, on the other hand, shows that memory utilization is now significantly lower than SM utilization, suggesting that the computation has become compute-bound rather than memory-bound:



We interpreted this as the memory bandwidth requirements of the kernel having been reduced, but the additional overhead, from computation, use of shared memory, and synchronization, being enough to decrease performance. Despite this, the results of profiling indicate that using shared memory will be useful in exposing opportunities to optimize computational resource use.

References

We followed the textbook's strategy for using shared memory in batched convolution.

2. Constant Memory

Since the kernel values are never modified, we can leverage constant memory to read values faster and further optimize the execution time.

How you identified the optimization opportunity: kernel memory is made up of constant values so they can be stored in constant memory to reduce load times while running the kernel

Why the approach is fruitful: It allows for faster load times when performing calculations using kernel memory

The effect: Overall, the time taken is decreased. This implementation improves timing across layers, making it a clear improvement.

With constant memory

```
Test Accuracy: 0.8714
-----
-          TIMINGS
-----
Layer 1 GPUPTime: 14.879225 ms
Layer 1 OpTime: 14.902041 ms
Layer 1 LayerTime: 605.238437 ms
Layer 2 GPUPTime: 54.966667 ms
Layer 2 OpTime: 55.004011 ms
Layer 2 LayerTime: 481.50531 ms
```

Base Implementation

```
Test Accuracy: 0.8714
-----
-          TIMINGS
-----
Layer 1 GPUPTime: 16.54464 ms
Layer 1 OpTime: 16.574783 ms
Layer 1 LayerTime: 650.947637 ms
Layer 2 GPUPTime: 58.609756 ms
Layer 2 OpTime: 58.639516 ms
Layer 2 LayerTime: 521.298612 ms
```

References: Lecture slides

Analysis using nsys:

Time(%)	Total Time	Calls	Average	Minimum	Maximum	Name
85.5	1284671642	6	214111940.3	16040	625229124	cudaMemcpy
13.3	199828379	8	24978547.4	80100	196121320	cudaMalloc
1.0	14442027	6	2407004.5	16596	14329357	cudaLaunchKernel
0.2	2824313	8	353039.1	69063	871088	cudaFree
0.0	168240	2	84120.0	83535	84705	cudaMemcpyToSymbol
0.0	20408	4	5102.0	2418	7891	cudaDeviceSynchronize
Generating CUDA Kernel Statistics...						
Generating CUDA Memory Operation Statistics...						
CUDA Kernel Statistics (nanoseconds)						
Time(%)	Total Time	Instances	Average	Minimum	Maximum	Name
100.0	270338390	2	135169195.0	43672815	226665575	conv_forward_kernel
0.0	2752	2	1376.0	1312	1440	do_not_remove_this_kernel
0.0	2656	2	1328.0	1312	1344	prefn_marker_kernel

This nsys output shows runtimes for the kernel with both constant memory and shared memory optimization. The first data set shows that even though we add a cudaMemcpyToSymbol call, the time it adds is worth the tradeoff as it significantly reduces the time other calls take. The second data set shows the total time our conv_forward_kernel takes which, again, is faster than the previous runs. This leads me to the conclusion that constant memory is worth implementing and speeds up the program's runtime.

3. Tuning with restrict and loop unrolling (considered as one optimization only if you do both)

We noticed that our for loops could be unrolled to obtain better performance results. This approach seemed fruitful because we could essentially half the number of iterations of the for loop and each iteration would do double the work. After implementing loop unrolling we found that the results were indeed fruitful:

After loop-unrolling:

Before loop-unrolling (w/ shared & constant mem):


```

Test Accuracy: 0.8714
-----
-                TIMINGS
-----
Layer 1 GPUSTime: 38.660198 ms
Layer 1 OpTime: 38.686886 ms
Layer 1 LayerTime: 685.301769 ms
Layer 2 GPUSTime: 202.451139 ms
Layer 2 OpTime: 202.481731 ms
Layer 2 LayerTime: 648.082145 ms

```

```

Test Accuracy: 0.8714
-----
-                TIMINGS
-----
Layer 1 GPUSTime: 43.350917 ms
Layer 1 OpTime: 43.374981 ms
Layer 1 LayerTime: 668.012451 ms
Layer 2 GPUSTime: 223.536424 ms
Layer 2 OpTime: 223.568904 ms
Layer 2 LayerTime: 691.185513 ms

```

We utilized the Nvidia documentation on unrolling loops [here](https://www.nvidia.com/docs/IO/116711/sc11-unrolling-parallel-loops.pdf).
<https://www.nvidia.com/docs/IO/116711/sc11-unrolling-parallel-loops.pdf>

Additionally, we used `__restrict__` to promise to the compiler that no pointer aliasing occurs. We pair this with using the keyword “const” as well to ensure that the input pointers are read only. After pairing this with loop-unrolling, we got these results which show further improvement than just loop-unrolling alone:

```

Test Accuracy: 0.8714
-----
-                TIMINGS
-----
Layer 1 GPUSTime: 36.651829 ms
Layer 1 OpTime: 36.674325 ms
Layer 1 LayerTime: 628.448989 ms
Layer 2 GPUSTime: 203.304852 ms
Layer 2 OpTime: 203.331796 ms
Layer 2 LayerTime: 642.932878 ms

```

We utilized the Nvidia documentation [here](https://developer.nvidia.com/blog/cuda-pro-tip-optimize-pointer-aliasing/) and [this Piazza post](https://pi Piazza post).
<https://developer.nvidia.com/blog/cuda-pro-tip-optimize-pointer-aliasing/>
<https://pi Piazza post>

According to the CUDA kernel statistics from nsys, our total time, average time, minimum time, and maximum time in `conv_forward_kernel` have all gone down as well.

Here are the nsys time statistics before restrict and loop-unrolling:

CUDA Kernel Statistics (nanoseconds)						
Time(%)	Total Time	Instances	Average	Minimum	Maximum	Name
100.0	274269946	2	137134973.0	43529320	230740626	conv_forward_kernel
0.0	2720	2	1360.0	1312	1408	prefn_marker_kernel
0.0	2624	2	1312.0	1248	1376	do_not_remove_this_kernel

Here are the nsys time statistics after restrict and loop-unrolling:

CUDA Kernel Statistics (nanoseconds)						
Time(%)	Total Time	Instances	Average	Minimum	Maximum	Name
100.0	230854050	2	115427025.0	38140081	192713969	conv_forward_kernel
0.0	2464	2	1232.0	1184	1280	prefn_marker_kernel
0.0	2464	2	1232.0	1216	1248	do_not_remove_this_kernel

Additional Optimizations

Based on our analyses, we built our combined kernel with the following configuration:

- No shared memory
- An additional guard to reduce computation in some cases
- Reduced branching when unrolling

Organization (PM4)

We had each team member implement one optimization, measuring and writing up the analysis for the results on a shared Google doc. The parts of the report for previous PMs was taken from Benoit and Steven's report.

The optimizations were assigned as follows:

- Benoit - Constant memory
- Steven - loop unrolling + restrict
- Leon - Shared memory (convolution)

Final Submission

1. FP16

FP16 can be used to make computations faster by allowing the processor to use less memory. Using the `__half2` data type, we can perform two operations in a single register.

How you identified the optimization opportunity: The GPU does many arithmetic operations so speeding those up could result in a significant faster runtime.

Why the approach is fruitful: It allows for faster arithmetic computation times in the GPU.

The effect: Overall, the total time taken is very similar and OpTime is shown to be consistently a bit slower. I believe this is because the time taken to transform the values into `__half2`'s is similar to the time gained by having more efficient arithmetic. (This run was made on top of previous constant memory optimization and has slight variations due to shared usage of rai GPU)

With FP16

```
Test Accuracy: 0.8716
-----
-              TIMINGS
-----
Layer 1 GPTime: 20.221661 ms
Layer 1 OpTime: 20.245853 ms
Layer 1 LayerTime: 619.425948 ms
Layer 2 GPTime: 72.630452 ms
Layer 2 OpTime: 72.661076 ms
Layer 2 LayerTime: 500.675067 ms
```

Base implementation

```
Test Accuracy: 0.8714
-----
-              TIMINGS
-----
Layer 1 GPTime: 17.439955 ms
Layer 1 OpTime: 17.445587 ms
Layer 1 LayerTime: 657.218768 ms
Layer 2 GPTime: 57.648491 ms
Layer 2 OpTime: 57.681419 ms
Layer 2 LayerTime: 531.559798 ms
```


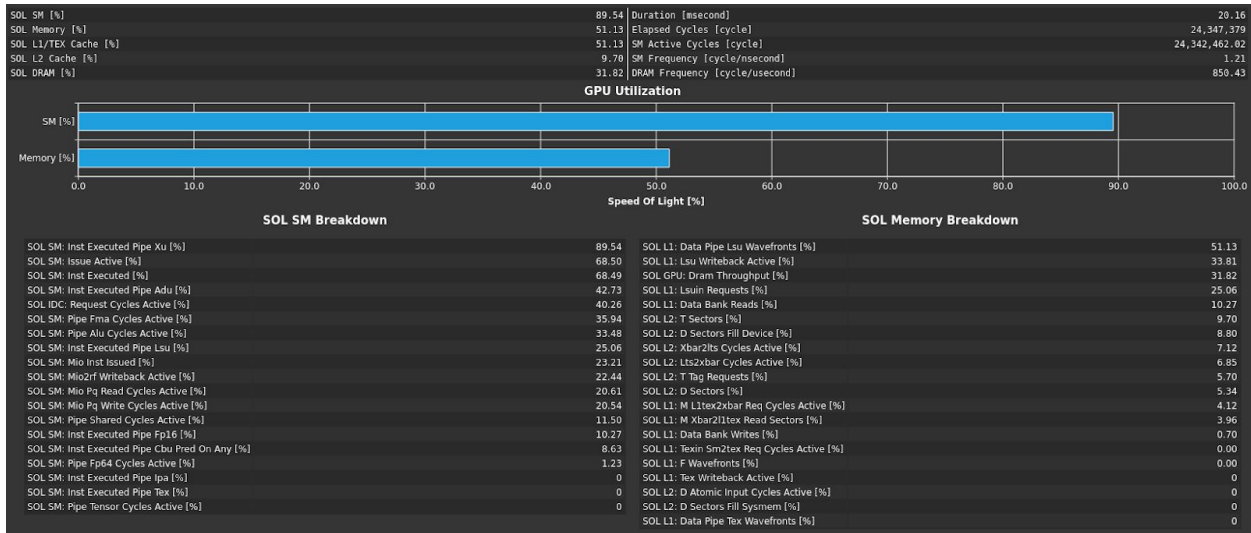
References:

https://docs.nvidia.com/cuda/cuda-math-api/group__CUDA__MATH__HALF__MISC.html

Nsys profiling:

Time(%)	Total Time	Calls	Average	Minimum	Maximum	Name
78.2	1154380961	6	192396826.8	12190	595379172	cudaMemcpy
19.9	293716529	8	36714566.1	65110	283262850	cudaMalloc
1.6	22973183	6	3828863.8	20383	22843141	cudaLaunchKernel
0.3	4302010	8	537751.3	62683	1901363	cudaFree
0.0	170611	2	85305.5	79983	90628	cudaMemcpyToSymbol
0.0	20382	4	5095.5	2346	10860	cudaDeviceSynchronize
Generating CUDA Kernel Statistics...						
Generating CUDA Memory Operation Statistics...						
CUDA Kernel Statistics (nanoseconds)						
Time(%)	Total Time	Instances	Average	Minimum	Maximum	Name
100.0	99316539	2	49658269.5	20211741	79104798	conv_forward_kernel
0.0	2848	2	1424.0	1344	1504	do_not_remove_this_kernel
0.0	2592	2	1296.0	1248	1344	prefn_marker_kernel
CUDA Memory Operation Statistics (nanoseconds)						
Time(%)	Total Time	Operations	Average	Minimum	Maximum	Name
93.2	975187745	2	487593872.5	400908765	574278980	[CUDA memcpy DtoH]
6.8	71072613	6	11845435.5	1440	37264362	[CUDA memcpy HtoD]
CUDA Memory Operation Statistics (KiB)						
	Total	Operations	Average	Minimum	Maximum	Name
	1722500.0	2	861250.0	722500.000	1000000.0	[CUDA memcpy DtoH]
	538919.0	6	89819.0	0.004	288906.0	[CUDA memcpy HtoD]
Generating Operating System Runtime API Statistics...						
Operating System Runtime API Statistics (nanoseconds)						
Time(%)	Total Time	Calls	Average	Minimum	Maximum	Name
33.4	96940081579	983	98616563.2	26805	102484943	sem_timedwait
33.1	96090883747	974	98655938.1	37483	100884043	poll
22.1	64025321315	2	32012660657.5	23809445709	40215875606	pthread_cond_wait
11.0	32008771817	64	500137059.6	500056944	500172939	pthread_cond_timedwait
0.3	725794388	101	7186083.0	3418	724679098	open64
0.1	216241321	866	249701.3	1151	103824375	ioctl
0.0	20531431	9072	2263.2	1123	18684	read
0.0	3653754	98	37283.2	1151	1686661	mmap
0.0	267227	5	53445.4	36316	60382	pthread_create
0.0	101327	18	5629.3	1262	15990	munmap
0.0	99071	24	4128.0	1089	14355	fopen
0.0	94821	15	6321.4	2539	14748	write
0.0	78130	3	26043.3	3444	43684	fopen64
0.0	76458	3	25486.0	11230	50078	fgets

Nsight profiling:



2. Sweeping various parameters to find best values (block sizes, amount of thread coarsening)

We had the idea of implementing this optimization from looking through the list of possible optimizations on the documentation. We expected that this approach would be effective because we haven't fiddled around with the block sizes at all up to this point, and thought doing so (in addition with thread coarsening) would yield a beneficial result.

Timings based on milestone 4 code:

```
Test Accuracy: 0.8714
-----
-                TIMINGS
-----
Layer 1 GPUSTime: 27.200195 ms
Layer 1 OpTime: 27.239074 ms
Layer 1 LayerTime: 686.241985 ms
Layer 2 GPUSTime: 116.891079 ms
Layer 2 OpTime: 116.917959 ms
Layer 2 LayerTime: 569.835636 ms
```

Timings based on milestone 4 code with thread coarsening and block size optimization:

```
Test Accuracy: 0.8714
-----
-                TIMINGS
-----
Layer 1 GPUSTime: 26.554858 ms
Layer 1 OpTime: 26.584009 ms
Layer 1 LayerTime: 638.380628 ms
Layer 2 GPUSTime: 63.119745 ms
Layer 2 OpTime: 63.143233 ms
Layer 2 LayerTime: 497.560514 ms
```

CUDA Kernel statistics before applying this optimization:

Generating CUDA Memory Operation Statistics...

CUDA Kernel Statistics (nanoseconds)

Time(%)	Total Time	Instances	Average	Minimum	Maximum	Name
100.0	147930892	2	73965446.0	27212334	120718558	conv_forward_kernel
0.0	2944	2	1472.0	1376	1568	do_not_remove_this_kernel
0.0	2688	2	1344.0	1312	1376	prefn_marker_kernel

CUDA Kernel statistics after applying this optimization:

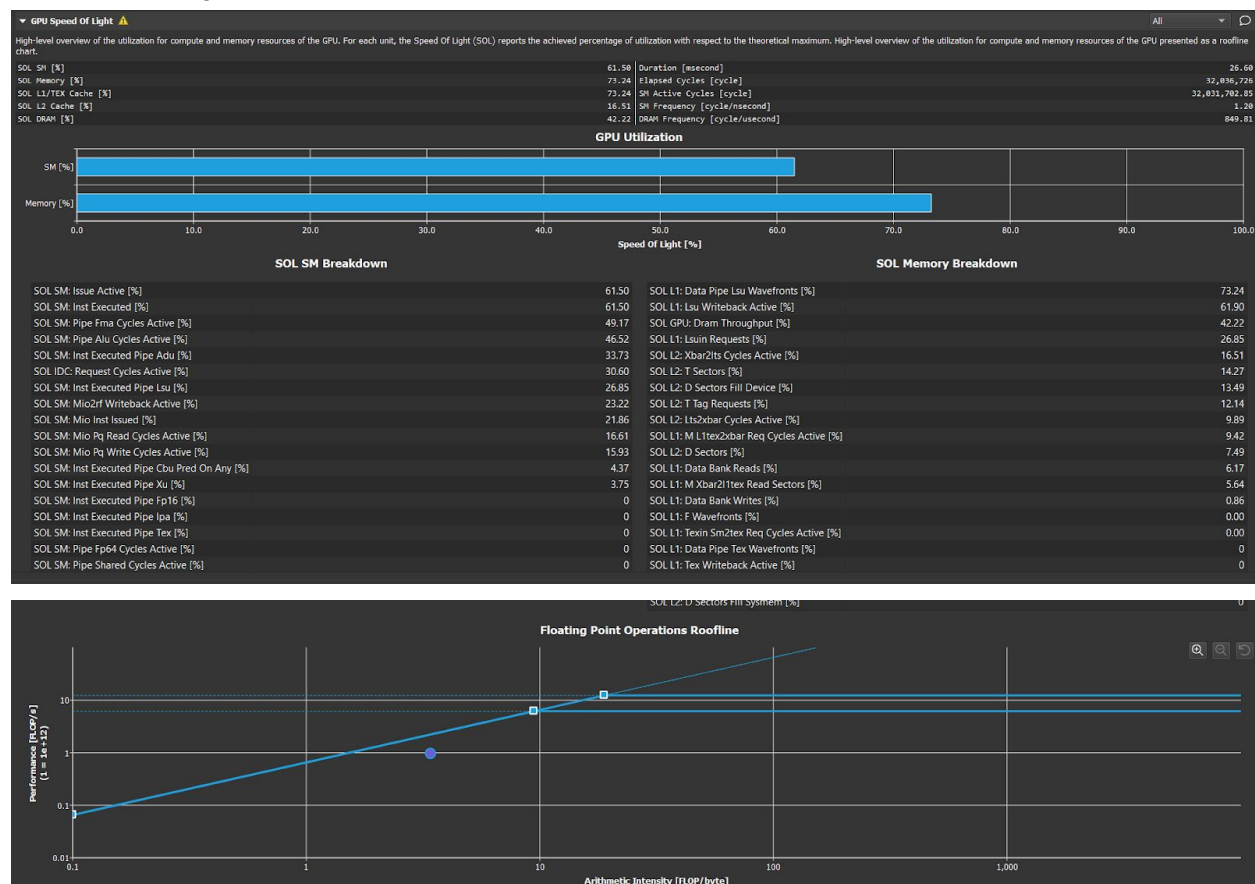
Generating CUDA Memory Operation Statistics...

CUDA Kernel Statistics (nanoseconds)

Time(%)	Total Time	Instances	Average	Minimum	Maximum	Name
100.0	90457753	2	45228876.5	23692975	66764778	conv_forward_kernel
0.0	5824	2	2912.0	2752	3072	prefn_marker_kernel
0.0	2816	2	1408.0	1376	1440	do_not_remove_this_kernel

As you can see, the total, average, and minimum time using our conv_forward_kernel are lower after applying this optimization to the milestone 4 implementation.

Here is the Nsight Compute for further reference:



External references: <https://dl.acm.org/doi/pdf/10.1145/3194242>

3. Two kernel implementations for different layer sizes

Motivation

While attempting to implement reduction across input channels within a block, we had to reduce the tiling width to remain under the thread limit for each block. We noticed that performance improved with smaller tiles, and measured the time for various block sizes against the base submission from PM4, taking the minimum Op times over three runs. The smaller layer, with one input channel, performed better with tiles of width 16, while the larger layer performed better with tiles of width 8. The obvious way to exploit this seemed to be to treat the single channel case as a special case, using a tile width of 16 and writing a kernel that eliminated the code required to handle multiple channels. Convolutions with multiple input channels used a tile width of 8.

Effect

The performance benefit appears to have been marginal over the base implementation with a tile width of 16. This was the timing output from one of our measurements:

```
-----  
-                TIMINGS  
-----  
  
Layer 1 GPUTime: 17.898761 ms  
Layer 1 OpTime: 17.923081 ms  
Layer 1 LayerTime: 617.52697 ms  
Layer 2 GPUTime: 55.093589 ms  
Layer 2 OpTime: 55.122356 ms  
Layer 2 LayerTime: 518.597423 ms
```

We also provide a table of timing measurements from the base implementation:

Block Size	Layer 1 Time (Minimum of 3)	Layer 2 Time (Minimum of 3)
8x4	30.99	65
8x8	23.48	51.61
16x16	17.92	55.90

Hybrid (This optimization)	17.92	55.12
-----------------------------------	--------------	-------

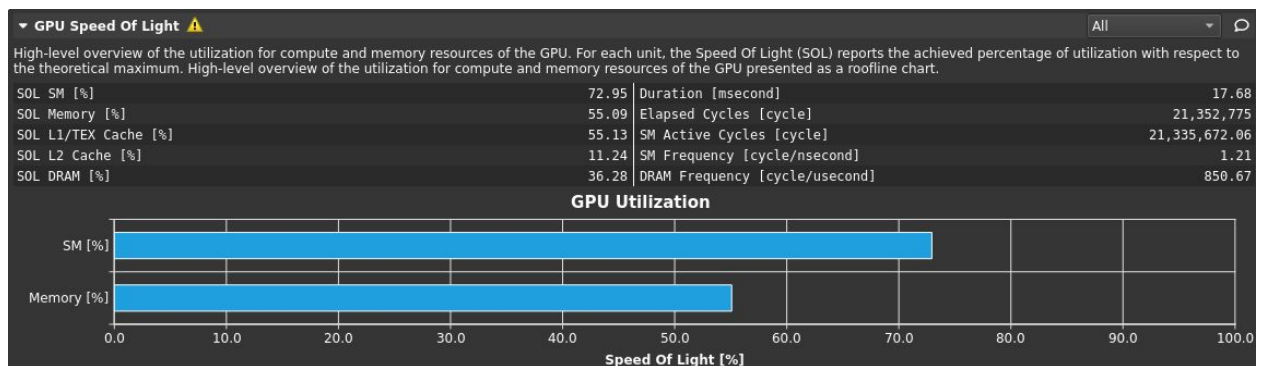
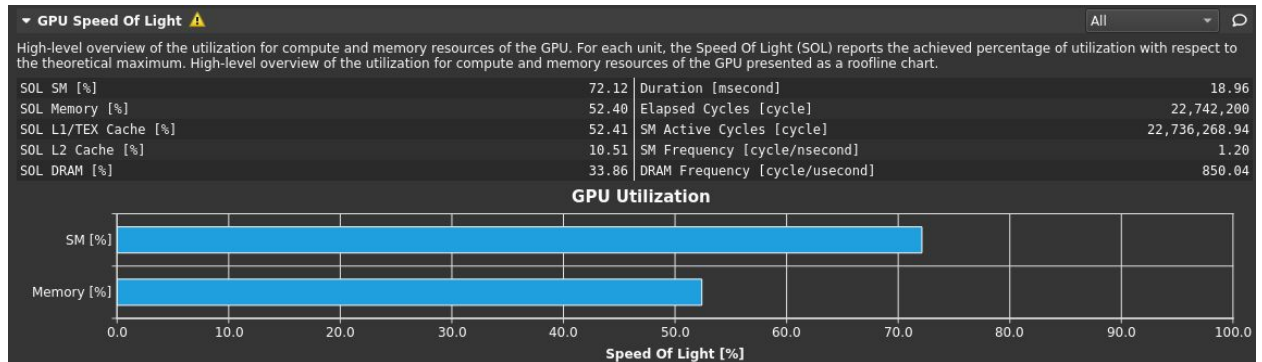
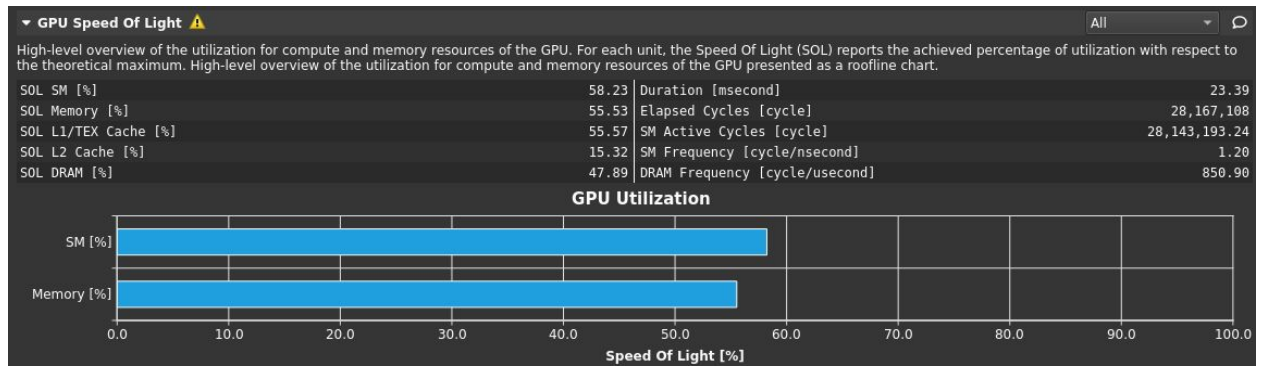
It is possible that the measurement from the 8x8 case was an outlier and that a tile width of 8 does not normally perform better for layer 2. Our two-kernel implementation, however, at least matches the best baseline performance, and has access to any benefit that a tile width of 8 has for the performance on multiple input channels.

Nsys Output


CUDA API Statistics (nanoseconds)						
Time(%)	Total Time	Calls	Average	Minimum	Maximum	Name
80.0	1113727705	6	185621284.2	16123	575747388	cudaMemcpy
18.7	260156502	8	32519562.7	58602	251424688	cudaMalloc
1.2	16080994	6	2680165.7	15805	15964591	cudaLaunchKernel
0.2	2261191	8	282648.9	54754	817778	cudaFree
0.0	349850	2	174925.0	172553	177297	cudaMemcpyToSymbol
0.0	17141	4	4285.2	2435	7462	cudaDeviceSynchronize
Generating CUDA Kernel Statistics...						
Generating CUDA Memory Operation Statistics...						
CUDA Kernel Statistics (nanoseconds)						
Time(%)	Total Time	Instances	Average	Minimum	Maximum	Name
75.4	55285311	1	55285311.0	55285311	55285311	conv_forward_kernel
24.6	18065741	1	18065741.0	18065741	18065741	conv_forward_kernel_onechannel
0.0	3872	2	1936.0	1312	2560	prefn_marker_kernel
0.0	2720	2	1360.0	1280	1440	do_not_remove_this_kernel
CUDA Memory Operation Statistics (nanoseconds)						
Time(%)	Total Time	Operations	Average	Minimum	Maximum	Name
91.3	945911286	2	472955643.0	389012847	556898439	[CUDA memcpy DtoH]
8.7	89667333	6	14944555.5	1472	48032463	[CUDA memcpy HtoD]
CUDA Memory Operation Statistics (KiB)						
	Total	Operations	Average	Minimum	Maximum	Name
	1722500.0	2	861250.0	722500.000	1000000.0	[CUDA memcpy DtoH]
	538919.0	6	89819.0	0.004	288906.0	[CUDA memcpy HtoD]

Analysis

Nsight was used to try to make sense of these results. We first looked at the speed of light statistics for the 8x8, 16x16, and hybrid implementations, on the first layer. The hybrid and 16x16 implementations had significantly higher SM utilization than with a block width of 8:

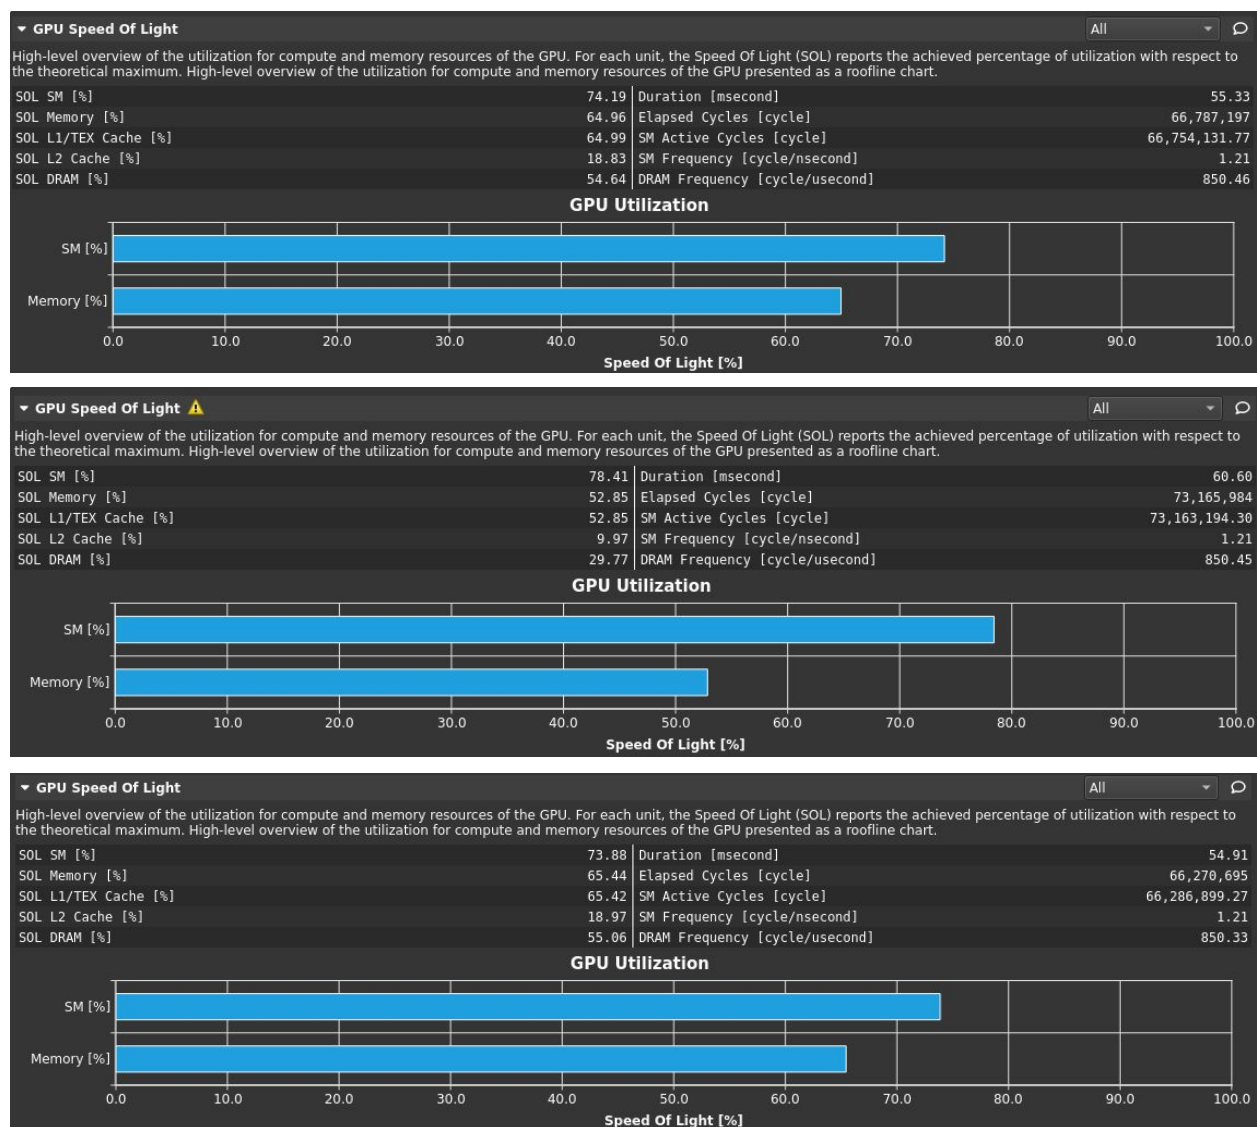


We further saw that that the 8x8 implementation stalled much more on scoreboard dependencies:



From these statistics we concluded that increasing the block width from 8 to 16 improved data reuse when accessing the L1 cache, which would have been used heavily because our base implementation does not use shared memory. This would not have affected invocations on the second layer as much, because it has data reuse across input channels as well as spatially in a particular input channel.

We had more trouble analyzing the performance behavior of the second kernel invocation, with the multi-channel input:



A higher memory utilization for 8x8 tiling is expected, as there is less data reuse within each tile. The occupancy statistics show that occupancy was a little higher for the kernel invocations that used 8x8 tiling, so it's possible that a higher block limit made more warps available to hide latency over:

► Occupancy				
Occupancy is the ratio of the number of active warps per multiprocessor to the maximum number of possible active warps. Another way to view occupancy is the percentage of the hardware's ability to process warps that is actively in use. Higher occupancy does not always result in higher performance, however, low occupancy always reduces the ability to hide latencies, resulting in overall performance degradation. Large discrepancies between the theoretical and the achieved occupancy during execution typically indicates highly imbalanced workloads.				
Theoretical Occupancy [%]	50	Block Limit Registers [block]		16
Theoretical Active Warps per SM [warp]	32	Block Limit Shared Mem [block]		32
Achieved Occupancy [%]	47.08	Block Limit Warps [block]		32
Achieved Active Warps Per SM [warp]	30.13	Block Limit SM [block]		32

► Occupancy				
Occupancy is the ratio of the number of active warps per multiprocessor to the maximum number of possible active warps. Another way to view occupancy is the percentage of the hardware's ability to process warps that is actively in use. Higher occupancy does not always result in higher performance, however, low occupancy always reduces the ability to hide latencies, resulting in overall performance degradation. Large discrepancies between the theoretical and the achieved occupancy during execution typically indicates highly imbalanced workloads.				
Theoretical Occupancy [%]	50	Block Limit Registers [block]		4
Theoretical Active Warps per SM [warp]	32	Block Limit Shared Mem [block]		32
Achieved Occupancy [%]	44.02	Block Limit Warps [block]		8
Achieved Active Warps Per SM [warp]	28.17	Block Limit SM [block]		32

► Occupancy				
Occupancy is the ratio of the number of active warps per multiprocessor to the maximum number of possible active warps. Another way to view occupancy is the percentage of the hardware's ability to process warps that is actively in use. Higher occupancy does not always result in higher performance, however, low occupancy always reduces the ability to hide latencies, resulting in overall performance degradation. Large discrepancies between the theoretical and the achieved occupancy during execution typically indicates highly imbalanced workloads.				
Theoretical Occupancy [%]	50	Block Limit Registers [block]		16
Theoretical Active Warps per SM [warp]	32	Block Limit Shared Mem [block]		32
Achieved Occupancy [%]	47.01	Block Limit Warps [block]		32
Achieved Active Warps Per SM [warp]	30.09	Block Limit SM [block]		32

References

N/A

Contributions (Final Milestone)

- Benoit - FP16
- Steven - Parameter sweeping / thread coarsening
- Leon - Two-kernel optimization.